

E-Cell AI CRM - System Report

1. Executive Summary


This system implements a five-module AI-native CRM for the fictional E-Cell company. It manages 520+ customers, 1,050+ tickets, and 2,500+ interaction events with LangChain summarization, LangGraph agent workflows, cohort retention analysis, and a live HEART framework dashboard backed by a production FastAPI API.

2. Architecture Decisions

2.1 Stack

Layer	Choice	Rationale
Backend	FastAPI + SQLite	Fast OpenAPI docs, zero external DB setup for demo
LLM (primary)	Ollama / Llama-3-8B (Option A)	On-prem privacy, no API cost
LLM (alternate)	Gemini 1.5 Flash (Option B)	Strong reasoning, free tier
LLM (fallback)	Rule-based mock	Ensures demo works offline
Agents	LangGraph state machine	Explicit routing → generate → escalation pipeline
Analytics	pandas + scikit-learn	Cohort metrics and churn classifier

2.2 Module Boundaries



Each module is independently runnable via `python -m src.<module>`.

3. LLM Selection Rationale

We support **dual-mode inference**:

1. **Ollama (Option A)** - Selected when `LLM_PROVIDER=ollama`. Suitable for evaluation environments requiring data residency.
2. **Gemini 1.5 Flash (Option B)** - Selected when `GEMINI_API_KEY` is set. Used for richer summarization during dataset generation.
3. **Mock provider** - Deterministic templates when neither backend is reachable; prevents demo failure.

Ticket summarization uses a dedicated system prompt requiring factual grounding. Agent responses pass through a **hallucination guard** that penalizes unsourced dollar amounts, missing ticket citations, and overconfident language.

4. Agent Prompt Design & LangGraph State Machine

4.1 State Schema

```
AgentState = {  
    customer_id, query, context, ticket_data,  
    route_category, priority, draft_response,  
    sources, confidence, escalated, hallucination_flags  
}
```

4.2 Graph Flow

Node	Function
`load_context`	Pulls short/long-term memory + open tickets
`route`	Classifies category (billing/technical/account/general)
`generate`	Drafts response with source citation
`escalation`	Flags low-confidence urgent cases for supervisor handoff

4.3 Summarization Chain

LangChain-style chain in `TicketSummarizationChain`:

- **Input:** ticket title, description, customer metadata
- **Output:** summary, `key_issues[]`, urgency, `suggested_response`
- **Configurable:** tone, `max_length`

5. Interaction Memory

Per-customer buffers in `customer_memory` table:

Layer	Capacity	Behavior
-------	----------	----------

Short-term	20 turns	Recent chat/ticket context
Long-term	Compressed summary	LLM-compressed when short-term overflows

Cross-session retrieval concatenates long-term summary + recent turns for agent queries.

6. Cohort Segmentation Methodology

6.1 Cohort ID Assignment

`cohort_id = {acquisition_month}_{industry}_{product_tier}`

Example: 2025-11_FinTech_Enterprise

Configurable dimensions: acquisition date, industry vertical, product tier, behavioral tags.

6.2 Retention Curve

For each cohort, we compute 6 monthly periods:

- **Active customer** at period *p*: $\text{engagement_score} \geq \max(20, 70 - 8p)$ AND $\text{tenure_days} \geq 30p$
- **Retention rate** = active / cohort_size

6.3 Churn Scoring

Heuristic score (production baseline):

$\text{churn_prob} = 0.5 * (1 - \text{engagement}/100) + 0.3 * \min(\text{tickets}/15, 1) + 0.2 * (1 - \text{tenure}/365)$

ML validation: Logistic regression on engagement, tickets, tenure. Labels use top-quartile heuristic scores (75th percentile cutoff) with stratified train/test split. Reports precision, recall, and F1 on held-out set.

6.4 Churn Model Evaluation

Metric	Source
Precision	sklearn on heuristic-labeled holdout
Recall	Same
F1	Harmonic mean
Coverage	% customers with cohort_id assigned

7. HEART Framework Metric Definitions

H - Happiness

Signal	Computation
CSAT	Mean `csat_score` on resolved tickets; falls back to mean engagement
NPS proxy	$\text{(promoters - detractors)} / N \times 100$ where promoter = engagement ≥ 70

E - Engagement

Signal	Computation
Active customers	Count with engagement ≥ 30
Ticket open rate	Open tickets / total tickets
Session depth proxy	Total interactions / customers

A - Adoption

Signal	Computation
AI-assisted rate	Tickets with `ai_assisted=true` / total
Onboarding completion	Customers with tenure ≥ 7 days
Memory adoption	Customers with memory buffer / total

R - Retention

Signal	Computation
Monthly retention	Customers with engagement ≥ 25 / cohort size
Churn flags	engagement < 35 AND tickets > 5
Avg lifespan	Mean tenure_days

T - Task Success

Signal	Computation
Resolution rate	Resolved+Closed / total tickets
FCR proxy	Non-escalated resolutions / total
Escalation rate	Escalated tickets / total
AI vs human ratio	AI-resolved / human-resolved

All metrics recompute from live SQLite data on each API call.

8. RBAC & Audit

Role	Permissions
Agent	CRUD customers/tickets, summarize, query agent
Supervisor	Agent + cohort read
Admin	Full access
Analytics	Read-only customers, tickets, cohorts, HEART

Every core endpoint returns `audit` block: `timestamp`, `agent_id`, `source_confidence`, `latency_ms`, `role`.

9. System Evaluation Metrics (Stage 6)

Category	Endpoint	Metric
HEART	<code>`/api/v1/heart`</code>	All 5 dimensions + per-metric breakdown
Agent Quality	<code>`/api/v1/query/agent`</code>	confidence, hallucination_flags
Cohort Accuracy	<code>`/api/v1/evaluation/metrics`</code>	churn model P/R/F1
Resolution Rate	<code>`/api/v1/evaluation/metrics`</code>	closure rate, AI ratio
Latency	All endpoints	<code>`audit.latency_ms`</code> per request

10. Dataset Specification

Entity	Count	Span
Customers	520	6 industries × 3 tiers
Tickets	1,050	billing, technical, account, general
Interactions	2,500	email, chat, call, ticket, portal
Time range	~6 months	Dec 2025 – Jun 2026

First 20 ticket descriptions are LLM-enhanced; remainder use realistic templates. Deduplication on email (customers) and ticket ID at ingest.

11. Known Limitations & Future Work

- SQLite suitable for demo; PostgreSQL recommended for production concurrency
- Churn labels derived from heuristic for ML eval until real churn events exist
- Latency benchmarks under 500 concurrent ops require separate load test (locust/k6)

- bcrypt hashing prepared but demo uses plaintext passwords for evaluator convenience

12. File Reference


File	Purpose
`src/crm.py`	Module 1
`src/agents.py`, `src/memory.py`	Module 2
`src/cohort.py`	Module 3
`src/heart.py`	Module 4
`api/app.py`	Module 5
`dashboard/index.html`	Live HEART UI
`SYSTEM_REPORT.md`	This document

Appendix A: Live Evaluation Metrics

Evaluation Results (Live CRM Data)

Metric	Value
Avg CSAT (Happiness)	3.73 / 5
Active customers (Engagement)	430
AI-assisted ticket rate (Adoption)	39.0%
Monthly retention (Retention)	88.3%
Ticket resolution rate (Task Success)	62.5%
Churn model F1	1.00
Churn-flagged customers	130
Largest cohort M+6 retention	70.0%

Appendix B: HEART Metrics Chart



Appendix C: Cohort Retention Curve

